

Software paper for submission to the Journal of Open Research Software

Please submit the completed paper to: editor.jors@ubiquitypress.com

(1) Overview

Title

spinsim: a GPU optimised solver of spin one and half quantum systems

Paper Authors

1. Tritt, Alex;
2. Morris, Joshua;
3. Hochstetter, Joel;
4. Anderson, R. P.;
5. Saunderson, James;
6. Turner, L. D.;

Paper Author Roles and Affiliations

1. School of Physics & Astronomy, Monash University, Victoria 3800, Australia.
Primary author of the released packages.
2. School of Physics & Astronomy, Monash University, Victoria 3800, Australia.
Present address: Faculty of Physics, University of Vienna, 1010 Vienna, Austria.
Author of first version of code.
3. School of Physics & Astronomy, Monash University, Victoria 3800, Australia.
Present address: School of Physics, University of Sydney, NSW 2006, Australia.
Optimization and extension to spin one of first version of code.
4. School of Molecular Sciences, La Trobe University, PO box 199, Bendigo, Victoria 3552, Australia.
Original conception of first version of code.
5. Department of Electrical and Computer Systems Engineering, Monash University, Victoria 3800, Australia.
Advice on numerical analysis.
6. School of Physics & Astronomy, Monash University, Victoria 3800, Australia.
Original conception of released version of algorithm.

Abstract

spinsim is a python package that simulates spin half and spin one quantum mechanical systems following a time dependent Shroedinger equation. It makes use of the *numba.cuda* llvm compiler [1], and other optimisations, to allow for fast and accurate evaluation on *Nvidia Cuda* [2] compatible systems using GPU parallelisation. *spinsim* is available for install on *PyPI*, and the source code is available on *github*. The initial use case for the package will be to simulate quantum sensing based Bose Einstein condensate (BEC) experiments for the Monash University School of Physics and Astronomy spinor BEC lab, but will be useful in simulating any range of spin one or half quantum systems with involved time dependent

Hamiltonians in fields of nuclear magnetic resonance (NMR), nuclear quadrupole resonance (NQR) and magnetic resonance imaging (MRI) experiments and quantum sensing, and with the spin one systems of nitrogen vacancy centres (NVCs) and BECs.

Keywords

Time dependent Schroedinger equation; Spin one; Spin half; Integrator; GPU; Solver; python; numba;

Introduction

Motivation

Ultracold atoms have proven their effectiveness in state of the art technologies in quantum metrology [3], the use of quantum mechanics to make precise measurements of small signals. Design of such measurement protocols requires many steps of verification, including simulation. This is especially important, since running real experiments can be expensive and time consuming, and thus it is more practical to debug such protocols quickly and cheaply on a computer. In general, any design of experiments using spin systems could benefit from a fast, accurate method of simulation.

In the past, the spinor Bose Einstein condensate (spinor BEC) lab at Monash University used an in-house, cython based script *AtomicPy* [4], which this package is originally based off, and standard differential equation solvers (such as *Mathematica*'s function *NDSolve*) to solve the Schroedinger equation for these spin one systems, which are sometimes approximated to spin half for a faster execution time. However, these methods are not completely optimised for our use case, and therefore come with some issues.

First of all, while the execution time for these solvers is acceptable for running a small number of experiments, for certain experiments involving large arrays of independent atom clouds (which require thousands of simulations to be run), this time accumulates to the order of many hours.

Secondly, the Schroedinger equation has the geometric property of being norm preserving. In other words, the time evolution operator for a system between two points in time must be unitary. This is important, because all quantum states live on the boundary of a ball of length one in their Hilbert space. For many numerical methods like those in the Runge Kutta family, the approximations used might not be norm preserving, and the evaluated quantum state may well diverge to infinite norm, or converge to zero if run for too many iterations.

Thirdly, our system (and similar spin systems) can be very oscillatory. In standard conditions for our application, the expected spin projection of a system can rotate in space (alternatively viewed as a point rotating around an abstract object known as a Bloch sphere) at a rate of 700kHz. Standard integration methods require very

small time steps in order to accurately depict these oscillations. We found that the integration method used by AtomicPy has an accuracy of only order 10^{-2} when the time step is set to the small value of 10ns.

Implementation and architecture

Mathematical methods

Parallelisation

In general, *spinsim* solves the Schroedinger equation,

$$\frac{d}{dt}\psi(t) = -iH(t)\psi(t),$$

with $i^2 = -1$, the state, $\psi(t)$ a time dependent, 3 (2) dimensional, unit complex vector, and $H(t)$ a time dependent 3 by 3 (2 by 2) complex matrix for spin one (spin half) systems. However, in *spinsim*, the Hamiltonian $H(t)$ is instead defined by its vector representation in Lie Algebra. For example, when set to spin half mode, the *spinsim* package solves time dependent Schroedinger Equations of the form

$$\frac{d}{dt}\psi(t) = -i2\pi(f_x(t)J_x + f_y(t)J_y + f_z(t)J_z)\psi(t),$$

where $i^2 = -1$, $\psi(t) \in \mathbb{C}^2$, and the spin half spin projection operators are given by

$$J_x = \frac{1}{2} \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}, \quad J_y = \frac{1}{2} \begin{pmatrix} 0 & -i \\ i & 0 \end{pmatrix}, \quad J_z = \frac{1}{2} \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}.$$

The source of the system is the collection of energy functions $f_x(t), f_y(t), f_z(t)$, with t in units of s and f in units of Hz that control the dynamics of the system. The user must define a function that returns a sample of these source functions when a sampling time is input. In physical terms, these functions would represent the x, y, z components of a magnetic field applied to an atom being simulated, for example.

Similarly, when *spinsim* is set to spin one mode, it can solve Schroedinger equations of the form

$$\frac{d}{dt}\psi(t) = -i2\pi(f_x(t)J_x + f_y(t)J_y + f_z(t)J_z + f_q(t)Q)\psi(t).$$

where now $\psi(t) \in \mathbb{C}^3$, and the spin one operators are given by

$$J_x = \frac{1}{\sqrt{2}} \begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{pmatrix}, \quad J_y = \frac{1}{\sqrt{2}} \begin{pmatrix} 0 & -i & 0 \\ i & 0 & -i \\ 0 & i & 0 \end{pmatrix},$$

$$J_z = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & -1 \end{pmatrix}, \quad Q = \frac{1}{3} \begin{pmatrix} 1 & 0 & 0 \\ 0 & -2 & 0 \\ 0 & 0 & 1 \end{pmatrix}.$$

J_x, J_y, J_z are regular spin operators, and Q is a quadratic operator, proportional to Q_{zz} as defined by Hamley et al [5], and Q_0 as defined by Di et al [6].

Given that $\psi(t)$ is a unit vector, it is possible to write $\psi(t)$ in terms of a unitary transformation $U(t, t_0)$ of the state $\psi(t_0)$, for any time t_0 . If the system is solved this way, using a geometric method, we can guarantee that $\psi(t)$ will always be a unit vector, which is not true with solvers in general.

$$\psi(t) = U(t, t_0)\psi(t_0)$$

This means that a time series for the state of the system can be evaluated by evaluating the time evolution operator between each of the sample times.

$$\begin{aligned} \psi_k &= U_k \psi_{k-1}, \text{ standing for} \\ \psi(t_k) &= U(t_k, t_{k-1})\psi(t_{k-1}), \text{ where} \\ t_k &= t_0 + Dt \cdot k \end{aligned}$$

with Dt the time step of the time series.

Also importantly, while each of the ψ_k must be evaluated sequentially, the value of the U_k is independent of the value of any ψ_{k_0} , or any other U_{k_0} . This means that the time evolution operators U_k can all be calculated in parallel, and it allows *spinsim* to use GPU parallelisation on the level of time sample points, so a speed up is achieved even if just a single simulation is run.

In summary, *spinsim* splits the time evolution of the full simulation into time evolution U_k within small time intervals $[t_{k-1}, t_k]$, which are each calculated massively in parallel on a GPU. When all the U_k are calculated CPU then multiplies the U_k together (a comparatively less demanding job than calculating them) to determine the ψ_k .

Rotating frame

If the rotating frame option is selected, the U_k are first calculated within a local rotating frame as U_k^r , which shrinks the size of the terms used in the calculation, increasing accuracy. The midpoint value $f_r = f_z(t_k + \frac{1}{2}Dt)$ is sampled to be removed from the source. This transforms the source functions as follows,

$$\begin{aligned} f_x^r(t) + if_y^r(t) &= e^{-i2\pi f_r t}(f_x(t) + if_y(t)), \\ f_z^r(t) &= f_z(t) - f_r, \\ f_q^r(t) &= f_q(t), \text{ for spin one.} \end{aligned}$$

This, (assuming that a midpoint sample is representative of an average value over the time interval), decreases the magnitude of $f_z(t)$, while leaving the other source

components at an equivalent magnitude. The rotation is then applied to obtain the lab frame time evolution operator U_k ,

$$\begin{aligned}
 U_k &= \exp(-i2\pi f_r J_z Dt) U_k^r, \\
 U_k &= \begin{pmatrix} e^{-i2\pi f_r Dt} & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & e^{i2\pi f_r Dt} \end{pmatrix} U_k^r, \text{ for spin one} \\
 U_k &= \begin{pmatrix} e^{-i\pi f_r Dt} & 0 \\ 0 & e^{i\pi f_r Dt} \end{pmatrix} U_k^r, \text{ for spin half.}
 \end{aligned}$$

Magnus based integration method

The integration method used in *spinsim* is the CF4 method from Auer et al [7]. Each of the U_k are split into products of time evolution operators between times separated by a smaller timestep,

$$\begin{aligned}
 U(t_k, t_{k-1}) &= U(t_k, t_k - dt) \cdots U(t_{k-1} + 2dt, t_{k-1} + dt) U(t_{k-1} + dt, t_{k-1}) \\
 U_k &= u_{L-1}^k \cdots u_0^k,
 \end{aligned}$$

with dt being the integration level, (ie, fine) time step.

The CF4 method is used to calculate each individual u_l^k . Let the fine sample time be given by $t_f = ldt + t_k$. Then as part of the CF4 method, the source functions are sampled at particular times based on the second order Gauss-Legendre quadrature,

$$\begin{aligned}
 f(t_1) &= (f_x(t_1), f_y(t_1), f_z(t_1), f_q(t_1)), \text{ with} \\
 t_1 &= t_f + \frac{1}{2}dt \left(1 - \frac{1}{\sqrt{3}}\right), \\
 f(t_2) &= (f_x(t_2), f_y(t_2), f_z(t_2), f_q(t_2)), \text{ with} \\
 t_2 &= t_f + \frac{1}{2}dt \left(1 + \frac{1}{\sqrt{3}}\right).
 \end{aligned}$$

The fine time evolution operator can then be calculated using

$$\begin{aligned}
 g_1 &= 2\pi dt \left(\frac{3 + 2\sqrt{3}}{12} f(t_1) + \frac{3 - 2\sqrt{3}}{12} f(t_2) \right) \\
 g_2 &= 2\pi dt \left(\frac{3 - 2\sqrt{3}}{12} f(t_1) + \frac{3 + 2\sqrt{3}}{12} f(t_2) \right) \\
 u &= \exp(-i(g_{2,x}J_x + g_{2,y}J_y + g_{2,z}J_z + g_{2,q}Q)) \\
 &\quad \cdot \exp(-i(g_{1,x}J_x + g_{1,y}J_y + g_{1,z}J_z + g_{1,q}Q)).
 \end{aligned}$$

Exponentiator

For all exponentiation, the exponentiator takes exponent in the form of

$$E(g) = \exp(-i(g_x J_x + g_y J_y + g_z J_z + g_q Q)), \text{ with} \\ g = (g_x, g_y, g_z, g_q)$$

For spin half, the default exponentiator is in an analytic form. For spin one, an exponentiator based on the Lie Trotter product formula [8] is used. Importantly, these two methods both use analytic forms of exponentials to construct the result, meaning that all calculated time evolution operators are unitary.

This also means that the package cannot solve arbitrary spin one quantum systems, as that would require the ability to exponentiate a point in the full, 8 dimensional Lie algebra of $\mathfrak{su}(3)$, rather than just the four dimensional subspace spanned by the subalgebra $\mathfrak{su}(2)$ spanned by $\{J_x, J_y, J_z\}$, and the single quadratic operator Q . Including the full algebra could be possible as a feature update if there is demand for it, though just including this subspace is sufficient for our application, and many others.

Software architecture

Integrator architecture

The integrator in the *spinsim* package calls a *numba.cuda.jit()*ed kernel to be run on a *Cuda* capable *Nvidia* GPU in parallel, with a different thread being allocated to each of the U_k . This returns when each of the U_k have been evaluated.

The thread starts by calculating t_k and, if the rotating frame is being used, f_r . The latter is done by sampling a (*numba.cuda.jit()*ed version of a) user provided python function f describing how to sample the source Hamiltonian. The code then loops over each fine time step dt to calculate the fine time evolution operators u_i^k .

Within the loop, the integrator enters a device function (ie a GPU subroutine, which is inline for speed) to sample $f(t)$, as well as calculate $e^{-i2\pi f_r t}$, at the sample times needed for the integration method. After this, it enters a second device function, which makes a rotating wave transformation as needed in a third device function, before calculating g values, and finally taking the matrix exponentiation in a fourth device function. u_i^k is premultiplied to U_k^r (which is initialised to 1), and the loop continues.

When the loop has finished, if the rotating frame is being used, U_k^r is transformed to U_k as above, and this is returned. Once all threads have executed, the state ψ_k is calculated in a (CPU) *numba.jit()*ed function from the U_k and an initial condition ψ_{init} .

Compilation of integrator

The *spinsim* integrator is constructed and compiled just in time, using *numba.cuda.jit()*. The particular device functions used are not predetermined, but are instead chosen based off user input to decide a closure. This structure has multiple advantages. Firstly, the source function f is provided by the user as a plain python function (that must be *numba.cuda.jit()* compatible). This allows users to define f in a way that compiles and executes fast, does not put many restrictions on the form of the function, and returns the accurate results of analytic functions (compared to the errors seen in interpolation). Compiling the simulator also allows the user to set metaparameters, and choose the features they want to use, in a way that does not require experience with the *numba.cuda* library. This was especially useful for running benchmarks comparing old integration methods to the new ones, like CF4. The default settings should be optimal for most users, although tuning the values of *Cuda* metaparameters *max_registers* and *threads_per_block* could improve performance for GPUs with a differing number of registers and *Cuda* cores to the mobile GTX1070 mainly used in testing here. Finally, just in time compilation also allows the user to select a target device other than *Cuda* for compilation, so the simulator can run, using the same algorithm, on a multicore CPU in parallel instead of a GPU, if the user so chooses.

This functionality is interfaced through an object of class *spinsim.Simulator*. The *Cuda* kernel is defined as per the user's instructions on construction of the instance, and it is used by calling the method *spinsim.Simulator.evaluate()*, which returns a results object including the time, state, time evolution operator, and expected spin projection (that is, Bloch vector. Note that the expected spin projection is calculated as a lazy parameter if needed, rather than returned by the simulator object).

Quality control Benchmarks

Speed

Benchmarks were performed using *sense.sim.benchmark*, by comparing evaluation speed of typical spin one metrology experiments on different devices. This is shown in Figure 1. The integration code was compiled by *numba* for single core CPUs, multicore CPUs, and *Nvidia Cuda*, and run on different devices. These devices are,

- Intel Core i7-8750H, a 6 core laptop processor released in 2018. Run with 16GiB of RAM. Air cooled (laptop fan). Base clock speed of 2.2GHz.
- AMD Ryzen 7 5800X, an 8 core desktop processor released in 2020. Run with 32GiB of RAM. Liquid cooled. Base clock speed of 3.8GHz.
- AMD Ryzen 9 5900X, a 12 core desktop processor released in 2020. Run with 32GiB of RAM. Air cooled. Base clock speed of 3.7GHz.

- Nvidia GeForce GTX 1070, a 2048 *Cuda* core laptop graphics processor released in 2016. Run with 8GiB of VRAM. Air cooled (laptop fan).
- Nvidia GeForce RTX 3070, a 5888 *Cuda* core desktop graphics processor released in 2020. Run with 8GiB of VRAM. Air cooled.
- Nvidia GeForce RTX 3080, an 8704 *Cuda* core desktop graphics processor released in 2020. Run with 10GiB of VRAM. Air cooled.

This benchmark shows the benefit to using parallelisation when solving this problem. Moving from a 6 core processor to a 12 core processor doubles the execution speed. Moving from a single core processor to a GPU increases performance by well over an order of magnitude. As an aside, liquid cooling allows the 8 core processor to increase its boost clock and outperform the 12 core processor.

Accuracy

Benchmarks were performed using *neural-sense.sim.benchmark* (where *neural-sense* [9] is the quantum sensing package that *spinsim* was written for), by comparing state evaluations of many different typical metrology experiments and finding the mean error introduced when the fine (integration) time step is increased.

Figure 2 shows the performance of *spinsim* when running in spin one mode. This shows that both using the Magnus based CF4 method and moving into a rotating frame give significant increases to accuracy. The HS (half step) method in the lab frame, with a time step of 10ns was the method used by *AtomicPy*, the previous code used by the group for simulating spin systems. Compared to this, the best performing *spinsim* method is 5 orders of magnitude more accurate, while executing in a time 2 orders of magnitude faster.

From Figure 3, one gets essentially the same accuracy for each method when working in spin half mode compared to spin one, if all else is kept constant.

Figure 4 shows that the Lie Trotter based exponentiator does limit the maximum accuracy obtainable, and for spin half systems, one can increase accuracy further (and decrease execution time) by using an analytic based exponentiator.

Comparison to alternatives

As mentioned before, compared to our previous code, *AtomicPy*, the best performing *spinsim* method is 5 orders of magnitude more accurate, while executing in a time 2 orders of magnitude faster. We had planned to benchmark against some popular generic solvers. One such solver was *qutip.sesolve()*, a solver in the popular quantum mechanics python library, *QuTip* [10]. However, due to a known bug with the library's dependencies, this was not installable on Windows 10, the operating system being used for testing, and so benchmarks for it could not be run. We also planned to benchmark accuracy against the generic solver *scipy.integrate.ivp_solve()* in the python library *SciPy* [11]. However, using the same *get_field* function as in




Figure 1: Evaluation speed of a typical spin one metrology experiment. Fine time step is 100ns, as determined to be ideal by the accuracy experiments. Experiments run for a duration of 100ms. Evaluation time is determined by an average of 100 similar experiments for each device.

the spin one benchmarks used for the *spinsim*, simulating with a fine time step of 500ns (the largest used during the *spinsim* benchmarks), we found that a single simulation ran in 153s, which is over three orders of magnitude slower than the most accurate *spinsim* simulations. This means that a generous projection for the




Figure 2: Fine time step benchmark for spin one systems. CF4 is the Magnus commutator free integrator, HS is the two sample exponential integrator used in *AtomicPy*, MP is a single sample exponential integrator, RF is use of the rotating frame, and LF is lab frame (no use of the rotating frame). HS and MP results are drawn on top of each other due to their similarities.

time it would take to run the same benchmarks that *spinsim* runs in 11 minutes, using *scipy* would be over six days; for an integration package not designed for this problem (and would therefore likely be less accurate).

Testing

The simulator as a whole has been functionally tested against well known analytic approximations of the behaviour spin systems. This was done for every combination of integrator settings possible when compiling the integrator. The system was benchmarked in terms of accuracy vs fine time step, again, using every possible combination of integrator settings. This confirms that no integrator diverges away from the limiting solution when the time step is decreased in magnitude. The Lie Trotter matrix exponentiator was tested separately from the full system, as well as benchmarked separately.

These tests and benchmarks were run as part of the *neural_sense* package. The simulator has also been used as part of the measurement protocol being developed




Figure 3: Fine time step benchmark for spin half systems, using the Lie Trotter based exponentiator. CF4 is the Magnus commutator free integrator, HS is the two sample exponential integrator used in *AtomicPy*, MP is a single sample exponential integrator, RF is use of the rotating frame, and LF is lab frame (no use of the rotating frame). HS and MP results are drawn on top of each other due to their similarities.

there, and it has been tested as part of those algorithms as well.

The kernel execution was profiled thoroughly, and changes were made to optimise VRAM and register usage and transfer. This was done specifically for the hardware of an Nvidia GTX1070, so one may get some performance increases by changing some GPU specific metaparameters when instantiating the *spinsim.Simulator* object.

A good way to confirm that *spinsim* is functioning properly after an installation is to run the tutorial code provided and compare the outputs. Otherwise, one can run the benchmarks and simulation protocols in *neural_sense.sim.benchmark*.

(2) Availability

Operating system

Developed on Windows 10. Tested on MacOS Big Sur.

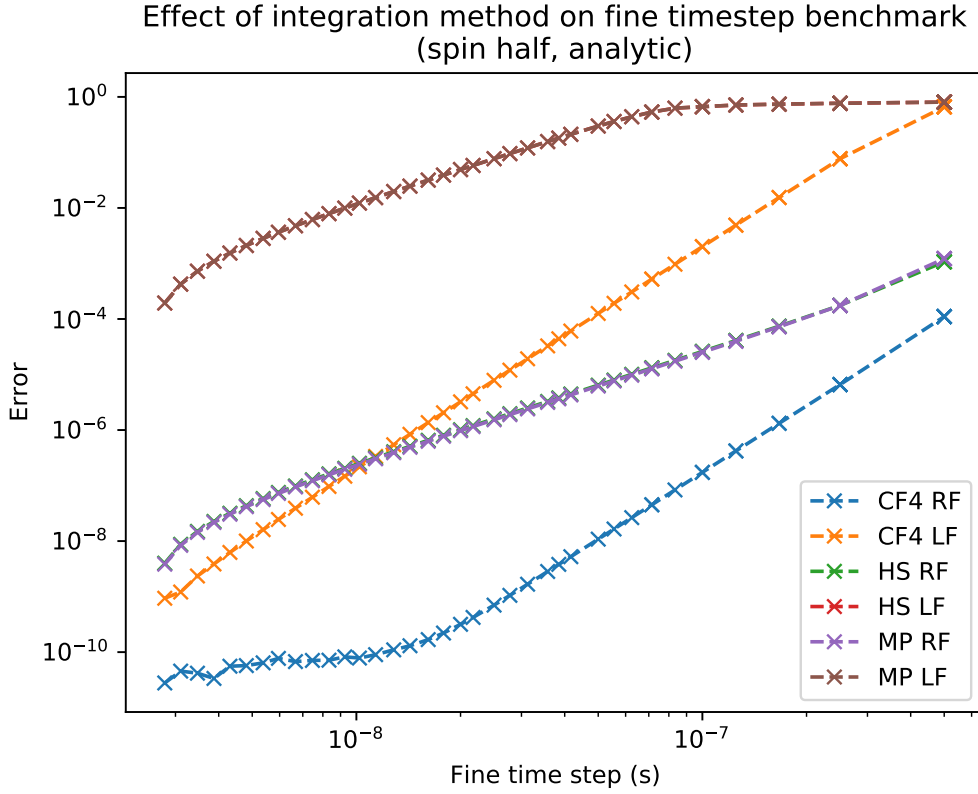


Figure 4: Fine time step benchmark for spin half systems, using the analytic based exponentiator. CF4 is the Magnus commutator free integrator, HS is the two sample exponential integrator used in *AtomicPy*, MP is a single sample exponential integrator, RF is use of the rotating frame, and LF is lab frame (no use of the rotating frame). HS and MP results are drawn on top of each other due to their similarities.

Programming language

Python (3.7 or greater)

Additional system requirements

To use the (default) *Nvidia Cuda* GPU parallelisation, one needs to have a *Cuda* compatible *Nvidia GPU*. For *Cuda* mode to function, one also needs to install the *Nvidia Cuda* toolkit. If *Cuda* is not available on the system, the simulator will automatically parallelise over multicore CPUs instead

Dependencies

numba (0.50.1 or greater)

numpy (1.19.3)

matplotlib (for example code, 3.2)

neuralsense (for benchmark code)

List of contributors

1. Alex Tritt

School of Physics & Astronomy, Monash University, Victoria 3800, Australia.

Primary author of the released packages.

2. Joshua Morris

School of Physics & Astronomy, Monash University, Victoria 3800, Australia.

Present address: Faculty of Physics, University of Vienna, 1010 Vienna, Austria.

Author of first version of code.

3. Joel Hockstetter

School of Physics & Astronomy, Monash University, Victoria 3800, Australia.

Present address: School of Physics, University of Sydney, NSW 2006, Australia.

Optimization and extension to spin one of first version of code.

4. Russell P. Anderson

School of Molecular Sciences, La Trobe University, PO box 199, Bendigo, Victoria 3552, Australia.

Original conception of first version of code.

5. James Saunderson

Department of Electrical and Computer Systems Engineering, Monash University, Victoria 3800, Australia.

Advice on numerical analysis.

6. Lincoln D. Turner

School of Physics & Astronomy, Monash University, Victoria 3800, Australia.

Original conception of released version of algorithm.

Software location:

Archive

Name: Monash Bridges

Persistent identifier: e.g. DOI, handle, PURL, etc.

Licence: BSD 3 Clause

Publisher: Alex Tritt

Version published: 1.0.0

Date published: dd/mm/yy

Code repository

Name: GitHub

Persistent identifier: <https://github.com/alexander-tritt-monash/spinsim>

Licence: BSD 3 Clause

Date published: 18/11/20

Language

English.

(3) Reuse potential

Use potential and limitations

spinsim will be useful for any research group needing quick, accurate, and / or large numbers of simulations involving spin half or spin one systems. This is immediately relevant to developing new quantum sensing protocols with spin one and spin half systems. The package will be used in the context of Bose Einstein Condensate (BEC) magnetic sensing protocol design by our lab, both within and outside the

project it was conceived for.

This project is to be able to measure neural signals using BECs. The electrical pulses made by neurons are currently measured using electrical probes, which is intrusive and damages the cells. We instead propose to sense the small magnetic fields that these electrical currents produce. Rubidium BECs can potentially be made sensitive enough to these tiny magnetic fields that they can be measured by them. An array of BECs can theoretically be made to sample the frequency components of the magnetic signal, essentially making a Fourier transform of it, which can be transformed back. Assuming that pulses do not occur very often (the signal is sparse in the time domain), a smaller BEC array with fewer frequency samples can be used to create a compressed measurement, reconstructed with theory from the field of compressive sensing [12]. *spinsim* was written to simulate possible measurement protocols for this, showing the behaviour of the array of spin one atoms interacting with the magnetic fields of the neurons, control signals, and noise.

Another example of spin based magnetic field sensing is the use of Nitrogen Vacancy Centres (NVCs). These are spin one structures found in diamond doped with Nitrogen atoms. This leaves a vacancy in a position adjacent to the Nitrogen atom, which pairs of electrons occupy to obtain the spin one properties. Similar to BECs, NVCs can be placed and addressed in 2D arrays in order to take many samples in one measurement. A paper was only recently released covering simulation experiments of magnetic neural pulse sensing using NVCs [13], which is something that *spinsim* could be useful for.

spinsim is designed to simulate small dimensional quantum systems, including large arrays of non-interacting spin systems. This means that it would not be able to integrate large arrays of entangled states or interacting particles. As a result, despite being fast at simulating qubits, it is inappropriate for the package to be used for quantum computing. In addition, *spinsim* is currently designed to integrate the time evolution of pure states only. This means that it may not be adequate for use in some Nuclear Magnetic Resonance (NMR) applications where relaxation [14] is important (or other kinds of simulations involving decoherence).

With these restrictions in mind, *spinsim* could be used for some simplified simulations in various areas of NMR. There are many atomic nuclei with spins of half (eg protons, Carbon 13) and, and fewer that have spins of one (eg Lithium 6, Nitrogen 14) [15], which, if relaxation and interactions between systems are not important for the application, *spinsim* could be used to simulate for spectroscopy experiments, for example. The inclusion of a quadrupole operator means that, with the same level of simplifications, *spinsim* should be able to simulate Nuclear Quadrupole Resonance (NQR) spectroscopy for spin one nuclei [16], such as Nitrogen 14, provided a suitable coordinate system is chosen. This technique measures energy level differences between levels split by electric field gradients, rather than static magnetic bias fields. Another possible use case could be for Magnetic Resonance Imaging (MRI) simulation and pulse sequence design. MRI uses measures the response of

spins of an array of spin half protons to a spatially varying pulse sequence [17], which essentially just corresponds to many separate *spinsim* simulations of spins at different positions in space. While this package offers some advantages over state of the art simulators in the field [18], with its use of quantum mechanics over classical mechanics, and its absence of rotating wave approximations, its parametrised pulse sequence definitions and geometric integrator, again, the lack of interacting particles and decoherence features are likely its downfall here.

Support

Documentation for *spinsim* is available on *Read the Docs*. This documentation contains a thorough tutorial on how to use the package, and installation instructions.

For direct support with the *spinsim* package, one can open an issue in the *github* repository. One can also use this contact to suggest extensions to the package. *spinsim* is planned to be maintained by the Monash University spinor BEC lab into the future.

Acknowledgements

Thank you to Hamish Taylor and Travis Hartley from the Monash University School of Physics and Astronomy spinor BEC lab, who have started using *spinsim* for their own projects and have given feedback of their user experience with the package.

Funding statement

If the software resulted from funded research please give the funder and grant number.

Competing interests

The authors declare that they have no competing interests.

References

- [1] Lam SK, Pitrou A, Seibert S. Numba: a LLVM-based Python JIT compiler. In: Proceedings of the Second Workshop on the LLVM Compiler Infrastructure in HPC - LLVM '15. Austin, Texas: ACM Press; 2015. p. 1–6. Available from: <http://dl.acm.org/citation.cfm?doid=2833157.2833162>.
- [2] Nickolls J, Buck I, Garland M, Skadron K. Scalable Parallel Programming with CUDA: Is CUDA the parallel programming model that application developers have been waiting for? Queue. 2008 Mar;6(2):40–53. Available from: <https://doi.org/10.1145/1365490.1365500>.
- [3] Huang J, Wu S, Zhong H, Lee C. Quantum Metrology with Cold Atoms. arXiv:13086092 [cond-mat, physics:physics, physics:quant-ph]. 2014 May;2:365–415. ArXiv: 1308.6092. Available from: <http://arxiv.org/abs/1308.6092>.
- [4] Morris J. QCmonk/Atomicpy; 2018. Original-date: 2018-08-03T03:07:58Z. Available from: <https://github.com/QCmonk/Atomicpy>.

- [5] Hamley CD, Gerving CS, Hoang TM, Bookjans EM, Chapman MS. Spin-nematic squeezed vacuum in a quantum gas. *Nature Physics*. 2012 Apr;8(4):305–308. Number: 4 Publisher: Nature Publishing Group. Available from: <https://www.nature.com/articles/nphys2245>.
- [6] Di Y, Wang Y, Wei H. Dipole–quadrupole decomposition of two coupled spin 1 systems. *Journal of Physics A: Mathematical and Theoretical*. 2010 Jan;43(6):065303. Publisher: IOP Publishing. Available from: <https://doi.org/10.1088%2F1751-8113%2F43%2F6%2F065303>.
- [7] Auer N, Einkemmer L, Kandolf P, Ostermann A. Magnus integrators on multi-core CPUs and GPUs. *Computer Physics Communications*. 2018 Jul;228:115–122. Available from: <http://www.sciencedirect.com/science/article/pii/S0010465518300584>.
- [8] Moler C, Van Loan C. Nineteen Dubious Ways to Compute the Exponential of a Matrix, Twenty-Five Years Later. *SIAM Review*. 2003;45(1):3–49. Publisher: Society for Industrial and Applied Mathematics. Available from: <https://www.jstor.org/stable/25054364>.
- [9] alexander-tritt monash. alexander-tritt-monash/neural-sense; 2020. Original-date: 2020-10-06T00:58:21Z. Available from: <https://github.com/alexander-tritt-monash/neural-sense>.
- [10] Johansson JR, Nation PD, Nori F. QuTiP 2: A Python framework for the dynamics of open quantum systems. *Computer Physics Communications*. 2013 Apr;184(4):1234–1240. Available from: <http://www.sciencedirect.com/science/article/pii/S0010465512003955>.
- [11] Virtanen P, Gommers R, Oliphant TE, Haberland M, Reddy T, Cournapeau D, et al. SciPy 1.0: fundamental algorithms for scientific computing in Python. *Nature Methods*. 2020 Mar;17(3):261–272. Number: 3 Publisher: Nature Publishing Group. Available from: <https://www.nature.com/articles/s41592-019-0686-2>.
- [12] Candes EJ, Wakin MB. An Introduction To Compressive Sampling. *IEEE Signal Processing Magazine*. 2008 Mar;25(2):21–30. Conference Name: IEEE Signal Processing Magazine.
- [13] Parashar M, Saha K, Bandyopadhyay S. Axon hillock currents enable single-neuron-resolved 3D reconstruction using diamond nitrogen-vacancy magnetometry. *Communications physics*. 2020;3(1):174.
- [14] Veshtort M, Griffin RG. SPINEVOLUTION: A powerful tool for the simulation of solid and liquid state NMR experiments. *Journal of Magnetic Resonance*. 2006 Feb;178(2):248–282. Available from: <http://www.sciencedirect.com/science/article/pii/S1090780705002442>.

- [15] Fuller GH. Nuclear Spins and Moments. *Journal of Physical and Chemical Reference Data*. 1976 Oct;5(4):835–1092. Publisher: American Institute of Physics. Available from: <https://aip.scitation.org/doi/abs/10.1063/1.555544>.
- [16] Bain AD, Khasawneh M. From NQR to NMR: The complete range of quadrupole interactions. *Concepts in Magnetic Resonance Part A*. 2004;22A(2):69–78. _eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1002/cmr.a.20013>. Available from: <https://onlinelibrary.wiley.com/doi/abs/10.1002/cmr.a.20013>.
- [17] McKinnon G. The Physics of Ultrafast MRI. In: Debatin JF, McKinnon GC, editors. *Ultrafast MRI: Techniques and Applications*. Berlin, Heidelberg: Springer; 1998. p. 1–51. Available from: https://doi.org/10.1007/978-3-642-80384-0_1.
- [18] Kose R, Setoi A, Kose K. A Fast GPU-optimized 3D MRI Simulator for Arbitrary k -space Sampling. *Magnetic Resonance in Medical Sciences*. 2019;18(3):208–218.

Copyright Notice

Authors who publish with this journal agree to the following terms:

Authors retain copyright and grant the journal right of first publication with the work simultaneously licensed under a Creative Commons Attribution License that allows others to share the work with an acknowledgement of the work's authorship and initial publication in this journal.

Authors are able to enter into separate, additional contractual arrangements for the non-exclusive distribution of the journal's published version of the work (e.g., post it to an institutional repository or publish it in a book), with an acknowledgement of its initial publication in this journal.

By submitting this paper you agree to the terms of this Copyright Notice, which will apply to this submission if and when it is published by this journal.